

# Ultra Low Latency File Synchronization Program Guide



Joshua Schaffer  
Trevor Wright  
Jayden Zebrowski  
Jazmyn Zurita

---

## Client Setup & Configuration

---

### Dependencies

The client is implemented in Rust and relies on eBPF (Extended Berkeley Packet Filter) for system-level monitoring. The following dependencies are required for successful compilation and operation:

- **Rust (Nightly)** – Required for advanced features and compatibility with the Aya eBPF framework and [cargo xtask](#).
- **BPF** – Enables loading and running of eBPF programs in the Linux kernel. Most modern distributions support this out of the box but must be enabled explicitly in some cases.
- **Clang** – Required for compiling the eBPF bytecode that runs in kernel space. Acts as the backend for [bpf-linker](#).
- **Cargo bpf-linker** – A Cargo plugin necessary for linking user-space and kernel-space components of the eBPF programs.
- **Linux Kernel Headers** – Provide the necessary definitions for building eBPF programs that interface correctly with the kernel. These must match the currently running kernel version.
- **Build-Essential** – A meta-package (or equivalent) that ensures basic build tools such as [make](#), [gcc](#), and [ld](#) are available for compiling system-level components. Mandatory on Debian-based systems.

Note: All dependencies should be installed and configured prior to attempting a build. See platform-specific setup instructions for further detail.

---

### Setup (Arch)

- `install rustup`
- `sudo pacman -S bpf`
- `sudo pacman -S linux-headers`
- `sudo pacman -S base-devel`
- `sudo pacman -S clang`
- `cargo install bpf-linker`
- `make install-aya-tool`

- make bindings
- 

### Setup (Debian)

- Sudo apt install linux-headers-\$(uname -r)
  - sudo apt install build-essential
  - sudo apt install clang
  - cargo install bpf-linker
  - make install-aya-tool
  - make bindings
  - Make run to run or cargo xtask run
- 

### Additional Setup

This section may not be required depending on your system configuration.

### Enable BPF LSM

If BPF LSM is not enabled, follow these steps to enable it. Reference: [AYA Book – LSM](#)

---

### Requirements

- **Kernel Version:** Must be at least 5.7
  - **LSM Check:** BPF must be listed as an enabled LSM
- 

### Verify BPF LSM

---

Run the following command:

```
cat /sys/kernel/security/lsm
```

Expected output should contain:

```
capability,lockdown,landlock,yama,apparmor,bpf
```

If **bpf** is missing from the list, BPF LSM must be enabled manually.

---

## Enable BPF LSM via GRUB

---

1. Edit the GRUB config:

```
sudo nano /etc/default/grub
```

2. Locate the line beginning with **GRUB\_CMDLINE\_LINUX** and append **,bpf** to the existing list of LSMs:

```
GRUB_CMDLINE_LINUX="lsm=[existing LSMs],bpf"
```

3. Rebuild GRUB using one of the following commands (depending on your distro):

```
sudo update-grub2
```

or

```
sudo grub2-mkconfig -o /boot/grub2/grub.cfg
```

or

```
sudo grub-mkconfig -o /boot/grub/grub.cfg
```

4. Reboot your system to apply changes:

```
sudo reboot
```

---

## Configs

Client configuration is controlled via the **config.json** file located in the project's root directory. Key configuration fields include:

- **Dns\_web\_addresses**: A list of server domain names or IP addresses to connect to. This list is dynamically re-read while the client is running, allowing updates without restart.
- **Watch\_dir**: The local directory to monitor for file changes. **Important**: The full path to this directory should be **shorter** than the server's path for consistency in file indexing and transmission.
- **Server\_port**: The port number used to connect to remote servers. By default, all connections use the same port, but this can be adjusted in code if multi-port support is desired.

- **Ignore\_ext: (Deprecated):** Previously used to ignore files based on extension, now no longer used.
- **Ignore:** A list of .gitignore-style rules. Files or directories matching any rule will be excluded from synchronization. Each rule should be a separate string in the configuration list.
- **File\_store\_time\_minutes:** Defines how long (in minutes) a file's full contents are kept in RAM before being purged. This controls RAM persistence for cached files.
- **Max\_total\_size\_mb:** Maximum total size (in MB) of all files held in RAM. If this limit is reached, new files will be skipped for caching and transmitted immediately without local storage.

---

## Code info

The client implementation is organized into multiple components. Below is a breakdown of the structure and functionality:

- Directory Structure
  - All user-facing code resides in the **ullfs** directory.
  - The **ullfs-ebpf** directory contains BPF-specific logic and may require modification if additional detections are needed.
  - All other code outside these directories should generally remain unchanged.
- Client\_tcp
  - This module is responsible for handling all data transmissions.
  - If enhancements such as compression or encryption are desired, the packet format should be adjusted within this module.

---

## ULLFS Protocol – Packet Format

The client sends data packets to the server following a custom protocol designed for performance and efficiency:

1. Filepath
  - a. Each character in the local file path is sent as a **u8**, relative to the **watch\_dir**.
2. Null Terminator
  - a. A single **0u8** byte is used to terminate the filepath string.
  - b. (File paths cannot contain a null byte, ensuring unambiguous termination.)
3. Flag (u8)
  - a. One-byte flag indicating the type of operation:
    - i. **1**: Full file send

- ii. 2: Delta send
- iii. 3: Delete file
- iv. 4: Move/Rename file
- v. 5: Create empty file
- vi. 6: Create directory

#### Additional Data by Flag

- Flag 1 – Full File Send
  - u64: File size in bytes
  - Raw file data
- Flag 2 – Delta Send
  - u64: Start byte (offset from beginning of file)
  - u64: End byte (offset from beginning of file)
  - u64: Data length
  - u64: Hash of the original file
  - Raw delta data
- Flags 3, 5, 6
  - No additional data beyond the flag.
- Flag 4 – Move/Rename File
  - New path as a string
  - Terminated with 0u8

---

### Packet Reception (Retransmission)

The client does not yet fully handle retransmission. If a retransmission packet is received from the server, it contains only a 0u8-terminated string (the file path). The client is expected to send back the full file, but this case is not yet implemented. The structure is in place for future support.

---

## Server Setup & Configuration

---

### Dependencies

The server is implemented in **Rust** and built on top of the **Steady-State Framework**. The following dependencies are required for successful compilation and execution:

- **Rust (Nightly)** – Required for compatibility with the Steady-State Framework macros and features.
- **Clang (optional)** – May be needed for certain underlying system dependencies.

- **Build-Essential** (*optional*) – Recommended for compiling native dependencies on Debian-based systems.

**Note:** In rare cases, the server may require additional dependencies that are typically needed by the client. However, these are not generally expected for standard operation.

---

## Configuration

Server configuration is controlled via the `config.json` file located in the project's root directory. Key configuration fields include:

- **Watch\_dir:** The directory to monitor for changes. Files within this directory are watched and changes are pushed to clients.
- **Server\_ip:** The server should always be bound to `0.0.0.0` to accept connections on all network interfaces.
- **Server\_port:** Set this to match the port expected by the client. Avoid using port `9100`, which is reserved by the Steady-State Framework for internal telemetry and visualization.

**Ignore:** This field is inherited from the client configuration and is currently unused in the server.

---

## Running the Server

To start the server:

```
cargo run
```

The application will compile and launch using the parameters defined in the `config.json`.

---

## Monitoring Telemetry

The Steady-State Framework provides real-time telemetry through a web interface. Once the server is running:

1. Open a web browser.
2. Navigate to: <http://127.0.0.1:9100>

This dashboard allows you to observe actor activity, channel throughput, and system utilization in real time.

---

## Codebase Overview

This server operates within the **Steady-State Framework**, which enforces a strict structural organization across the codebase. The primary logic for actor and channel creation resides in `/src/main.rs`. Developers should begin here when working on the core system configuration.

## Actor and Channel Configuration

Channel and actor creation follow global rules defined in `graph.channel_builder()` and `graph.actor_builder()`. These configurations apply uniformly unless explicitly overridden. The foundational method `base_channel_builder()` is used to instantiate a channel, taking two parameters: the transmitter and the receiver.

- To configure a channel's capacity, use the `.with_capacity()` method, which sets the maximum number of messages that can be in-flight at any given time.
- Finalize and instantiate the channel using the `.build()` method.

Actor instantiation is handled via `base_actor_builder.with_name()`, where you can define the actor's name, specify the logical core it runs on, and attach its communication channels.

For channel-specific or actor-specific behavior, refer to the [Steady-State crate documentation](#) for implementation guidelines and customization options.

## Actor Implementation

The behavior of each actor is defined in its corresponding source file located under `/src/actor/{actor_name}.rs`. These implementations follow a consistent structure based on the **FizzBuzz Steady-State** template.

The entry point for each actor is the `run()` function. Its parameters must align with those declared in the main file; as additional channels are connected, these parameters must be updated accordingly.

Telemetry is integrated via the `into_monitor!()` macro, which relays runtime statistics to the graph. The macro takes three arguments:

```
into_monitor!(context, [receiving_channels], [transmission_channels]);
```



**Important:** Rust's compiler and Rust Analyzer will not catch ordering mismatches in these macro arguments. If telemetry fails to output, this is a key area to inspect for potential issues.

The core logic of each actor is implemented in the `internal_behavior()` function. This function has the same parameters as `run()`, with an additional one for telemetry monitoring.

- Begin by locking all connected channels.
- Enter a `while` loop that processes until all transmission channels are **marked closed** and all receiving channels are **closed and empty**.
- Use `try_take(channel_name)` to consume messages from a receiving channel.
- Use `send_async(channel_name, message, SendSaturation)` to send messages through transmission channels.
- Call `relay_stats()` to update telemetry metrics, provided the corresponding attribute is defined in `graph.actor_builder()`.

## Modular Design

To simplify the codebase and improve maintainability, the server's processing logic has been modularized:

- All logic related to the **ULLFS protocol** is implemented in `/src/actor/handle_client.rs`.
- Configuration values, such as `watch_dir`, are centralized in the `config.json` file located in the root directory. This file acts as a single source of truth for critical runtime settings.
- Configuration parsing and accessors are implemented in `/src/actor/file_filter.rs`, allowing for easy instantiation and retrieval of configuration values using getter methods.

This modular and clearly structured approach ensures scalability and maintainability as the project grows.

---

---

## eBPF Components

---

1. Kernel-space eBPF program (`ullfs-ebpf/src/main.rs`): Implements LSM and FExit hooks
  2. User-space application (`ullfs/src/actor/ebpf_listener.rs`): Loads eBPF, processes events
- 

## Kernel-Space Implementation

---

### eBPF Constraints

- No standard library (`#![no_std]`)
- Limited stack space (512 bytes)
- Limited instruction count (1000000 instructions)
- No access to arbitrary memory
- No dynamic memory allocation
- Limited error handling

### Hook Types

1. LSM Hooks:
  - Primary source of filesystem events
  - Hooks: `path_unlink`, `inode_create`, `path_rename`, etc.
  - Used for capturing operations at security checkpoints
2. FExit Hook:
  - Currently only `vfs_write` is used
  - Captures write operations not visible through LSM hooks

### Kernel Structure Navigation

1. Dentry-centric design:
  - Path is only used to get to dentry: `path → dentry`
  - File is only used to get to path to get to dentry: `file → f_path → dentry`
  - All operations work with dentries as the primary data structure
2. Parameter extraction functions:
  - `try_arg_path`: Extracts dentry from path parameter
  - `try_arg_file`: Extracts dentry from file parameter
  - `try_arg_dentry`: Uses dentry parameter directly
  - `try_arg_dentry_parent`: Gets parent dentry when target is empty
3. Empty dentry issue:
  - When dentries are empty/incomplete at interception time
  - Only parent dentries contain valid info
  - Requires delayed inode verification in user space
4. Key kernel structures:
  - `dentry`: Directory entries forming filesystem hierarchy
  - `inode`: Unique file identifiers with metadata
  - `path`: File path representations (only used to access dentries)
  - `qstr`: Kernel string implementation for names

- All accessed safely via `bpf_probe_read_kernel`

## Data Structures

EventData:

```
struct EventData {  
    inod: u64,      // Primary inode number  
    inod2: u64,     // Secondary inode (for rename operations)  
    len: u16,       // Length of first path string  
    len2: u16,      // Length of second path string (for rename)  
    event_type: u8, // Type code for the event  
    rename_state: u8, // Rename state (different directory scenarios)  
}
```

BPF Maps:

- `INODEDATA`: Stores watched directory inode
- `PROGDATA`: Contains process information
- `BUF` and `BUFTWO`: Per-CPU arrays for path storage
  - Per-CPU design prevents synchronization issues
  - Each CPU gets dedicated buffer space
- `EVENTS`: Perf event array for user space communication

## Path Extraction

1. Path building process:
  - `pathToMap`: Traverses directory entries to build path
  - `dnameToMap`: Extracts name components from dentries
  - Path components stored in per-CPU arrays
  - Special character (!) used for empty components
2. Directory context:
  - `in_dir`: Traverses up directory tree
  - Checks if operations occur in watched directories
  - Compares against watched inode from `INODEDATA`
  - Filters events to watched directory only
3. Limitations:
  - Kernel paths stored in reverse order (leaf to root)
  - Limited buffer sizes require chunking
  - Special handling for rename/move operations

## Event Communication

1. Event details packaged in `EventData` structure
2. Path information stored in per-CPU arrays
3. Event notification sent to user space via perf event

---

## User-Space Implementation

---

## eBPF Program Management

1. Loads compiled bytecode using Aya framework
2. Attaches 20+ LSM hooks and the `vfs_write` FExit hook
3. Initializes maps:
  - Watched directory inode in `INODEDATA`
  - Process ID in `PROGDATA`
4. Sets up per-CPU event processing

## Two-Phase Event Processing

1. Initial phase (in eBPF):
  - Capture event type and available info
  - Extract inode numbers (often from parent dentry)
  - Store partial/full path information
2. Verification phase (in user space):
  - `check_inode`: Set when additional verification needed
  - Uses inode to find actual file/directory
  - Scans directory for matching inode
  - Reconstructs complete path from found file
  - Necessary because target dentries are often empty during syscall interception

## Path Reconstruction

1. `extract_filename` function:
  - Reads components from per-CPU arrays
  - Handles placeholder characters
  - Reverses path (kernel stores leaf to root)
2. Path components joined with proper separators
3. Relative paths calculated from watched directory

## Operation Type Resolution

1. Event type from kernel (0-31)
2. Context from inode verification
3. States (`mode` variable):
  - Mode 1: Rename operation
  - Mode 2: File creation
  - Mode 3: Directory creation
  - Mode 4: Deletion
  - Mode 5: Move operation
  - Mode 6: Move into watched directory
  - Mode 7: Creation of unknown type

## Special Cases

1. Rename operations:
  - `rename_state` field tracks directory context:
    - State 1: Both in watched dir

- State 2: Destination in watched dir (move to)
  - State 3: Source in watched dir (move from)
  - Requires tracking two paths simultaneously
- 2. Cross-boundary moves:
  - Files moved into watched area trigger full content send
  - Files moved out treated as deletes

### Final Processing

1. Event categorized by type
2. **TcpData** structure created with:
  - Operation type
  - File path(s)
3. Data sent through channel to other components

---

## Troubleshooting

---

### Client

---

**Issue:** `cargo xtask run` does not work

**Solution:** Try running the command using explicit binary targeting:

```
cargo run --bin xtask run
```

This occasionally happens after toolchain or crate updates.

---

### Server

---

**Issue:** Client fails to connect to the server

#### Possible Causes & Solutions:

- **Mismatched Ports:** Ensure that the client and server are configured to use the same port.
  - **Network Restrictions:** If running across different networks, make sure the target port is open and properly forwarded on the server's router.
  - **Debugging by Reversion:** If unsure what broke the connection, revert recent changes and gradually reapply them. Even minor changes can cause communication issues.
-